

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
import numpy as np
import seaborn as sns
from sklearn.preprocessing import LabelEncoder
from matplotlib import pyplot as plt
%matplotlib inline
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import ExtraTreesClassifier
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.feature_selection import mutual_info_classif
from sklearn.model_selection import cross_val_score
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix

### Importing all the libraries for the ananlysis
```

```
In [2]: train_data = pd.read_csv('healthcare-dataset-stroke-data.csv') # importing dataset
```

```
In [3]: train_data.columns
```

```
Out[3]: Index(['id', 'gender', 'age', 'hypertension', 'heart_disease', 'ever_married',
              'work_type', 'Residence_type', 'avg_glucose_level', 'bmi',
              'smoking_status', 'stroke'],
              dtype='object')
```

```
In [4]: train_data.dtypes #checking datatype for columns in the dataset
```

```
Out[4]: id                int64
gender                object
age                  float64
hypertension          int64
heart_disease          int64
ever_married          object
work_type              object
Residence_type        object
avg_glucose_level     float64
bmi                   float64
smoking_status        object
stroke                int64
dtype: object
```

```
In [5]: train_data.isnull().sum()      #checking if we have any null values in the dataset
```

```
Out[5]: id                0
gender                0
age                  0
hypertension         0
heart_disease        0
ever_married         0
work_type            0
Residence_type       0
avg_glucose_level    0
bmi                 201
smoking_status       0
stroke              0
dtype: int64
```

```
In [6]: train_data
```

```
Out[6]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	Male	67.0	0	1	Yes	Private	Urban
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural
2	31112	Male	80.0	0	1	Yes	Private	Rural
3	60182	Female	49.0	0	0	Yes	Private	Urban
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural
...
5105	18234	Female	80.0	1	0	Yes	Private	Urban
5106	44873	Female	81.0	0	0	Yes	Self-employed	Urban
5107	19723	Female	35.0	0	0	Yes	Self-employed	Rural
5108	37544	Male	51.0	0	0	Yes	Private	Rural
5109	44679	Female	44.0	0	0	Yes	Govt_job	Urban

5110 rows × 12 columns



```
In [7]: train_data['stroke'].value_counts(normalize=True)      #counting the value of stroke
```

```
Out[7]: 0    0.951272
1    0.048728
Name: stroke, dtype: float64
```

```
In [8]: train_data['bmi'].round() #Rounding the values of Bmi column so that we can conv
```

```
Out[8]: 0      37.0
        1      NaN
        2      32.0
        3      34.0
        4      24.0
        ...
       5105     NaN
       5106     40.0
       5107     31.0
       5108     26.0
       5109     26.0
       Name: bmi, Length: 5110, dtype: float64
```

```
In [9]: train_data.describe() # this shows the mathematical ground for each feature
```

```
Out[9]:
```

	id	age	hypertension	heart_disease	avg_glucose_level	bmi
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000

Checking for Null values

Manipulating the null value

```
In [10]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     5110 non-null   int64
1   gender                 5110 non-null   object
2   age                    5110 non-null   float64
3   hypertension           5110 non-null   int64
4   heart_disease          5110 non-null   int64
5   ever_married           5110 non-null   object
6   work_type              5110 non-null   object
7   Residence_type         5110 non-null   object
8   avg_glucose_level      5110 non-null   float64
9   bmi                    4909 non-null   float64
10  smoking_status         5110 non-null   object
11  stroke                 5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

```
In [11]: train_data.isnull().sum()      # checking total number of null values in the dataset
```

```
Out[11]: id                     0
gender                 0
age                    0
hypertension           0
heart_disease          0
ever_married           0
work_type              0
Residence_type         0
avg_glucose_level      0
bmi                    201
smoking_status         0
stroke                 0
dtype: int64
```

```
In [12]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                    5110 non-null   int64
1   gender                5110 non-null   object
2   age                   5110 non-null   float64
3   hypertension          5110 non-null   int64
4   heart_disease         5110 non-null   int64
5   ever_married          5110 non-null   object
6   work_type              5110 non-null   object
7   Residence_type        5110 non-null   object
8   avg_glucose_level     5110 non-null   float64
9   bmi                   4909 non-null   float64
10  smoking_status        5110 non-null   object
11  stroke                 5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

```
In [13]: #Changing Bmi null value to mean values of Bmi as for health purpose Average Bmi
train_data['bmi'] = train_data['bmi'].fillna(train_data['bmi'].mean())
```

```
In [14]: train_data
```

```
Out[14]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	Male	67.0	0	1	Yes	Private	Urban
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural
2	31112	Male	80.0	0	1	Yes	Private	Rural
3	60182	Female	49.0	0	0	Yes	Private	Urban
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural
...
5105	18234	Female	80.0	1	0	Yes	Private	Urban
5106	44873	Female	81.0	0	0	Yes	Self-employed	Urban
5107	19723	Female	35.0	0	0	Yes	Self-employed	Rural
5108	37544	Male	51.0	0	0	Yes	Private	Rural
5109	44679	Female	44.0	0	0	Yes	Govt_job	Urban

5110 rows × 12 columns



```
In [15]: train_data.isnull().sum()
```

```
Out[15]: id                0
gender              0
age                0
hypertension        0
heart_disease        0
ever_married         0
work_type            0
Residence_type       0
avg_glucose_level    0
bmi                 0
smoking_status       0
stroke              0
dtype: int64
```

In [16]: `train_data.describe()` *#after changing the null values we can see there was a diff*

Out[16]:

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237	
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.698018	
min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000	
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.800000	
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.400000	
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	32.800000	
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000	



In [17]: `train_data['bmi'] = train_data['bmi'].round()
#Bmi = train_data['bmi'].astype(int)`

In [18]: `train_data['bmi']`

Out[18]:

```
0      37.0
1      29.0
2      32.0
3      34.0
4      24.0
...
5105   29.0
5106   40.0
5107   31.0
5108   26.0
5109   26.0
Name: bmi, Length: 5110, dtype: float64
```

In [19]: `Bmi = train_data['bmi']`

```
In [20]: Bmi
```

```
Out[20]: 0      37.0  
         1      29.0  
         2      32.0  
         3      34.0  
         4      24.0  
         ...  
        5105     29.0  
        5106     40.0  
        5107     31.0  
        5108     26.0  
        5109     26.0  
        Name: bmi, Length: 5110, dtype: float64
```

```
In [21]: train_data['stroke'].value_counts()
```


```
Out[21]: 0      4861  
         1       249  
        Name: stroke, dtype: int64
```

```
In [22]: train_data['stroke'].value_counts(normalize=True)
```

```
Out[22]: 0      0.951272  
         1      0.048728  
        Name: stroke, dtype: float64
```

```
In [23]: train_data['stroke'].value_counts().plot.bar(xlabel='stroke') #Plotting bar to c
```

```
Out[23]: <AxesSubplot:xlabel='stroke'>
```





```
In [24]: train_data['ever_married'].value_counts()
```

```
Out[24]: Yes      3353  
        No       1757  
        Name: ever_married, dtype: int64
```

```
In [25]: train_data['ever_married'].value_counts().plot(kind='bar', color=['b','g']) #sin
```


```
Out[25]: <AxesSubplot:>
```



Plotting and Analysis using seaborn


```
In [26]: sns.countplot(train_data['hypertension']) #here we check the number of people w
```

```
Out[26]: <AxesSubplot:xlabel='hypertension', ylabel='count'>
```




```
In [27]: sns.catplot(x='stroke',y='age', data=train_data) #This graph shows number strok
```

```
Out[27]: <seaborn.axisgrid.FacetGrid at 0x24d0e640400>
```




```
In [28]: sns.countplot(train_data['heart_disease'])
```

```
Out[28]: <AxesSubplot:xlabel='heart_disease', ylabel='count'>
```




```
In [29]: sns.countplot(train_data['ever_married'])
```

```
Out[29]: <AxesSubplot:xlabel='ever_married', ylabel='count'>
```




```
In [30]: sns.countplot(train_data['bmi']) #Numerical Variable of bmi is higher for most of
```

```
Out[30]: <AxesSubplot:xlabel='bmi', ylabel='count'>
```




```
In [31]: sns.countplot(train_data['Residence_type']) #This graph shows there are equal num
```

```
Out[31]: <AxesSubplot:xlabel='Residence_type', ylabel='count'>
```




```
In [32]: sns.countplot(train_data['gender']) #this graph the dataset is more bended tow
```

```
Out[32]: <AxesSubplot:xlabel='gender', ylabel='count'>
```





```
In [33]: sns.countplot(train_data['avg_glucose_level']) #Numerical variable showing glucose
```

```
Out[33]: <AxesSubplot:xlabel='avg_glucose_level', ylabel='count'>
```



```
In [34]: sns.distplot(train_data['avg_glucose_level']) #Glucose level can impact on the outcome
plt.show()
train_data['avg_glucose_level'].plot.box(figsize=(16,5))
plt.show()
```



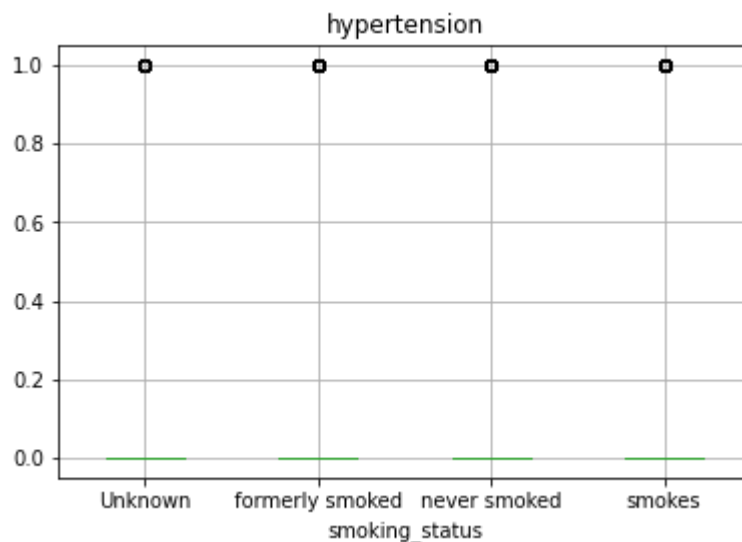
```
In [35]: lot(column='heart_disease', by = 'smoking_status') #comparing the boxplot in terms of heart disease
```

```
Out[35]: Text(0.5, 0.98, '')
```




```
In [36]: train_data.boxplot(column='hypertension', by = 'smoking_status')
plt.suptitle("")
```

```
Out[36]: Text(0.5, 0.98, '')
```




Bivariate Analysis

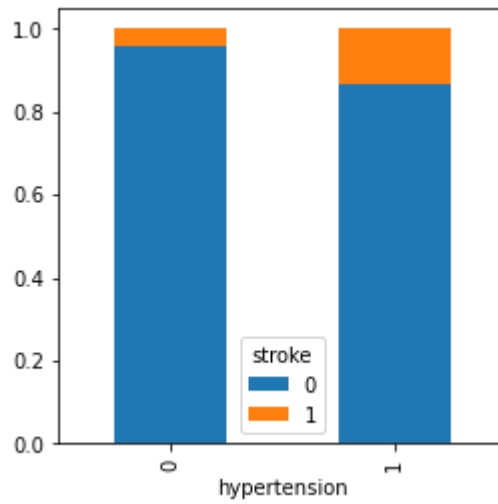
```
In [37]: Gender=pd.crosstab(train_data['gender'],train_data['stroke']) # Doing bivariate c
Gender.div(Gender.sum(1).astype(float), axis=0).plot(kind="bar",stacked=True,figs
plt.show()
```



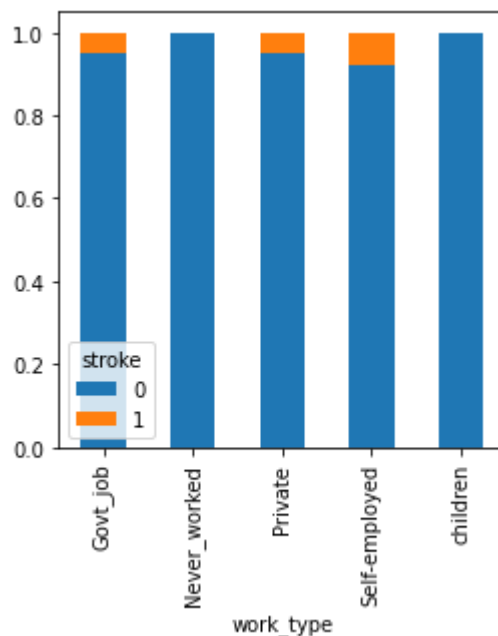
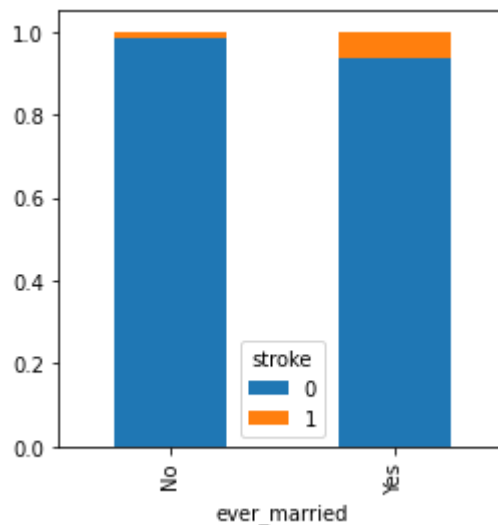
```
In [38]: heart=pd.crosstab(train_data['heart_disease'],train_data['stroke']) # Doing bivar
heart.div(heart.sum(1).astype(float), axis=0).plot(kind="bar",stacked=True,figsiz
plt.show()
```

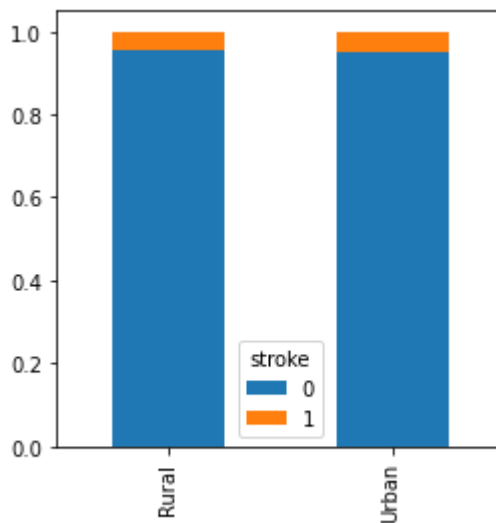



```
In [39]: b(train_data['hypertension'],train_data['stroke']) # Doing bivariate analysis with  
         um(1).astype(float), axis=0).plot(kind="bar",stacked=True,figsize=(4,4))
```



```
In [40]: #this is a Bivariate graph for remaining dependent features  
Married=pd.crosstab(train_data['ever_married'],train_data['stroke'])  
Working=pd.crosstab(train_data['work_type'],train_data['stroke'])  
Residence=pd.crosstab(train_data['Residence_type'],train_data['stroke'])  
Married.div(Married.sum(1).astype(float), axis=0).plot(kind="bar",stacked=True,figsize=(10,5))  
plt.show()  
Working.div(Working.sum(1).astype(float), axis=0).plot(kind="bar",stacked=True,figsize=(10,5))  
plt.show()  
Residence.div(Residence.sum(1).astype(float), axis=0).plot(kind="bar",stacked=True,figsize=(10,5))  
plt.show()
```





One Hot Encoding

```
In [41]: Work_type = pd.get_dummies(train_data['work_type'])      # Separating the variable
Residence_area = pd.get_dummies(train_data['Residence_type'])
Smoking_status = pd.get_dummies(train_data['smoking_status'])
Gender = pd.get_dummies(train_data['gender'])
Married = pd.get_dummies(train_data['ever_married'])
```

```
In [42]: Work_type
```

Out[42]:

	Govt_job	Never_worked	Private	Self-employed	children
0	0	0	1	0	0
1	0	0	0	1	0
2	0	0	1	0	0
3	0	0	1	0	0
4	0	0	0	1	0
...
5105	0	0	1	0	0
5106	0	0	0	1	0
5107	0	0	0	1	0
5108	0	0	1	0	0
5109	1	0	0	0	0

5110 rows × 5 columns

In [43]: Residence_area

Out[43]:

	Rural	Urban
0	0	1
1	1	0
2	1	0
3	0	1
4	1	0
...
5105	0	1
5106	0	1
5107	1	0
5108	1	0
5109	0	1

5110 rows × 2 columns

In [44]: Smoking_status

Out[44]:

	Unknown	formerly smoked	never smoked	smokes
0	0	1	0	0
1	0	0	1	0
2	0	0	1	0
3	0	0	0	1
4	0	0	1	0
...
5105	0	0	1	0
5106	0	0	1	0
5107	0	0	1	0
5108	0	1	0	0
5109	1	0	0	0

5110 rows × 4 columns

In [45]: Married

Out[45]:

	No	Yes
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1
...
5105	0	1
5106	0	1
5107	0	1
5108	0	1
5109	0	1

5110 rows × 2 columns

In [46]: Gender

Out[46]:

	Female	Male	Other
0	0	1	0
1	1	0	0
2	0	1	0
3	1	0	0
4	1	0	0
...
5105	1	0	0
5106	1	0	0
5107	1	0	0
5108	0	1	0
5109	1	0	0

5110 rows × 3 columns

Making New Variable seperate from Dataset

Making new dataset for every column so that we can't interfere with the main dataset

```
In [47]: Id = train_data['id']
```

```
In [48]: Hypertension = train_data['hypertension']
```

```
In [49]: Heart_disease = train_data['heart_disease']
```

```
In [50]: Age = train_data['age']
```

```
In [51]: Avg_glucose_level = train_data['avg_glucose_level']
```

```
In [52]: train_data.head()
```

Out[52]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	av
0	9046	Male	67.0	0	1	Yes	Private	Urban	
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	
2	31112	Male	80.0	0	1	Yes	Private	Rural	
3	60182	Female	49.0	0	0	Yes	Private	Urban	
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	

```
In [53]: Stroke = train_data['stroke']
```

```
In [54]: train_data['age'].value_counts()
```

Out[54]:

78.00	102
57.00	95
52.00	90
54.00	87
51.00	86
...	
0.48	3
1.40	3
0.16	3
0.08	2
0.40	2

Name: age, Length: 104, dtype: int64

```
In [55]: train_data['work_type'].value_counts()
```

```
Out[55]: Private      2925  
Self-employed    819  
children        687  
Govt_job        657  
Never_worked     22  
Name: work_type, dtype: int64
```

```
In [56]: train_data['Residence_type'].value_counts()
```

```
Out[56]: Urban      2596  
Rural      2514  
Name: Residence_type, dtype: int64
```

```
In [57]: train_data['avg_glucose_level'].value_counts()
```

```
Out[57]: 93.88      6  
72.49      5  
84.10      5  
91.68      5  
83.16      5  
..  
95.02      1  
120.09     1  
197.58     1  
99.91      1  
60.50      1  
Name: avg_glucose_level, Length: 3979, dtype: int64
```

```
In [58]: train_data['smoking_status'].value_counts()
```

```
Out[58]: never smoked    1892  
Unknown              1544  
formerly smoked      885  
smokes               789  
Name: smoking_status, dtype: int64
```

```
In [59]: Stroke.value_counts()
```


```
Out[59]: 0      4861  
1       249  
Name: stroke, dtype: int64
```

```
In [60]: finaldata = pd.concat([Gender, Age, Hypertension, Heart_disease, Married, Work_ty  
X = finaldata.iloc[:,0:-1] # concatenating the newly created and manipulated colu
```

Correlation Matrix

```
In [61]: # this shows the important feature required for the analysis as we can see every
corr=X.corr()
plt.figure(figsize=(15,10))
sns.heatmap(corr, annot = True, cmap="BuPu") # visualising the correlation matrix
```

Out[61]: <AxesSubplot:>



```
In [62]: y = finaldata['stroke']
```

```
In [63]: kbest = SelectKBest(score_func = chi2, k = 'all') # target number of features is
ordered_features = kbest.fit(X,y)
```

```
In [64]: train_data_scores = pd.DataFrame(ordered_features.scores_, columns=['Score']) # s
```

```
In [65]: train_data_columns = pd.DataFrame(X.columns, columns = ['Feature_name']) # save t
```

```
In [66]: feature_rank = pd.concat([train_data_scores,train_data_columns],axis=1) # combine
```


In [67]: `feature_rank.nlargest(12, 'Score') # rank the features by score - the scores are b`

Out[67]:

	Score	Feature_name
3	3635.226911	age
15	1718.285446	avg_glucose_level
5	87.987436	heart_disease
4	75.449498	hypertension
6	39.355836	No
12	31.111620	children
7	20.622787	Yes
18	17.607359	formerly smoked
11	16.584252	Self-employed
16	15.879229	bmi
17	11.139767	Unknown
9	1.126929	Never_worked

Extra Trees Classifier

In [68]: `# initialise the method`

```
model = ExtraTreesClassifier()
model.fit(X,y) # pass the input and output data to the method
```

Out[68]: `ExtraTreesClassifier()`


In [69]: `model.feature_importances_ # compute the scores`

Out[69]: `array([1.50852334e-02, 1.57031462e-02, 3.16749387e-07, 2.75320681e-01, 3.09127149e-02, 2.56070155e-02, 7.89691909e-03, 9.16796271e-03, 1.23054746e-02, 1.85154504e-05, 1.60240539e-02, 1.40040020e-02, 1.21828391e-03, 1.64421777e-02, 1.59429318e-02, 2.73874067e-01, 2.09193574e-01, 1.28708345e-02, 1.46831578e-02, 1.81327448e-02, 1.55961921e-02])`

In [70]: `ranked_features = pd.Series(model.feature_importances_, index = X.columns) # put`

```
In [71]: ranked_features.nlargest(11).plot(kind='bar') # rank and plot the scores
```

Out[71]: <AxesSubplot:>



Mutual Information Gain

```
In [72]: mu_ifo = mutual_info_classif(X,y)
```

```
In [73]: mu_data = pd.Series(mu_ifo, index = X.columns)
mu_data.sort_values(ascending=False)
```

```
Out[73]: age                0.036125
Yes                0.008051
avg_glucose_level  0.006257
heart_disease      0.006204
hypertension       0.005150
children           0.005073
bmi                0.004936
No                 0.004766
Urban              0.004236
Other              0.001745
Govt_job           0.001648
Female             0.001321
Never_worked       0.000705
Private            0.000494
formerly smoked    0.000453
Rural              0.000000
Self-employed      0.000000
Male               0.000000
Unknown            0.000000
never smoked       0.000000
smokes             0.000000
dtype: float64
```

```
In [74]: x_train,x_test,y_train,y_test = train_test_split(X,y, test_size = 0.2, random_state=42)
```

```
In [75]: def classify(model, x, y):
x_train,x_test,y_train,y_test = train_test_split(X,y, test_size = 0.2, random_state=42)
model.fit(x_train,y_train)
print('Accuracy is: ', model.score(x_test,y_test)*100)
score = cross_val_score(model,x,y,cv=5) # we set the number of folds to 5
print('Cross validation Accuaracy: ', np.mean(score)*100)
```

```
In [76]: LG_model = LogisticRegression() #making logistic regression model
classify(LG_model, X,y)
```


```
Accuracy is: 93.9334637964775
Cross validation Accuaracy: 95.12720156555773
```

```
In [77]: y_pred = LG_model.predict(x_test)
```

```
In [78]: cm = confusion_matrix(y_test,y_pred)
```

```
In [79]: sns.heatmap(cm,annot=True)
```

```
Out[79]: <AxesSubplot:>
```



```
In [80]: knn_model = KNeighborsClassifier(n_neighbors=5)
classify(knn_model,X,y)
```

Accuracy is: 93.73776908023484


Cross validation Accuaracy: 94.28571428571428

```
In [81]: y_pred = knn_model.predict(x_test)
```

```
In [82]: cm = confusion_matrix(y_test,y_pred)
```

```
In [83]: sns.heatmap(cm,annot=True)
```

```
Out[83]: <AxesSubplot:>
```



```
In [84]: SVC_model = SVC(kernel='linear', C = 1)  
classify(SVC_model, X,y)
```

Accuracy is: 93.9334637964775

Cross validation Accuaracy: 95.12720156555773

```
In [85]: SVC_model = SVC(kernel='linear', C = 1)
```

```
In [86]: SVC_model.fit(x_train,y_train)
```


```
Out[86]: SVC(C=1, kernel='linear')
```

```
In [87]: y_pred = SVC_model.predict(x_test) # save predictions in y_pred
```

```
In [88]: cm = confusion_matrix(y_test,y_pred)
```

```
In [89]: sns.heatmap(cm,annot=True) # plots the confusion matrix
```

```
Out[89]: <AxesSubplot:>
```



Trying the same process using Label Encoding

In [90]: train_data

Out[90]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	Male	67.0	0	1	Yes	Private	Urban
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural
2	31112	Male	80.0	0	1	Yes	Private	Rural
3	60182	Female	49.0	0	0	Yes	Private	Urban
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural
...
5105	18234	Female	80.0	1	0	Yes	Private	Urban
5106	44873	Female	81.0	0	0	Yes	Self-employed	Urban
5107	19723	Female	35.0	0	0	Yes	Self-employed	Rural
5108	37544	Male	51.0	0	0	Yes	Private	Rural
5109	44679	Female	44.0	0	0	Yes	Govt_job	Urban

5110 rows × 12 columns



```
In [91]: cols = ['gender', 'ever_married', 'work_type', 'Residence_type', 'smoking_status']
le = LabelEncoder() # initialising the necessary function t
for col in cols:
    train_data[col] = le.fit_transform(train_data[col])
```

```
In [92]: train_data
```

```
Out[92]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	1	67.0	0	1	1	2	1
1	51676	0	61.0	0	0	1	3	0
2	31112	1	80.0	0	1	1	2	0
3	60182	0	49.0	0	0	1	2	1
4	1665	0	79.0	1	0	1	3	0
...
5105	18234	0	80.0	1	0	1	2	1
5106	44873	0	81.0	0	0	1	3	1
5107	19723	0	35.0	0	0	1	3	0
5108	37544	1	51.0	0	0	1	2	0
5109	44679	0	44.0	0	0	1	0	1

5110 rows × 12 columns




```
In [93]: corr=train_data.corr()
```



```
In [94]: plt.figure(figsize=(15,10))
sns.heatmap(corr, annot = True, cmap="BuPu")
```

Out[94]: <AxesSubplot:>



```
In [95]: A = train_data.drop(['stroke'],axis=1)
        b = train_data['stroke']
```

```
In [96]: kbest = SelectKBest(score_func = chi2, k = 'all') # target number of features is
        ordered_features = kbest.fit(A,b)
```

```
In [97]: df_scores = pd.DataFrame(ordered_features.scores_, columns=['Score'])
```

```
In [98]: df_columns = pd.DataFrame(X.columns, columns = ['Feature_name'])
```

```
In [99]: feature_rank = pd.concat([df_scores,df_columns],axis=1)
```

```
In [100]: feature_rank.nlargest(12,'Score')
```

Out[100]:

	Score	Feature_name
2	3635.226911	Other
0	2556.735918	Female
8	1718.285446	Govt_job
4	87.987436	hypertension
3	75.449498	age
5	20.622787	heart_disease
9	15.879229	Never_worked
10	3.369423	Private
6	2.925901	No
7	0.600717	Yes
1	0.239001	Male

```
In [101]: model = ExtraTreesClassifier()
        model.fit(A,b)
```

Out[101]: ExtraTreesClassifier()


```
In [102]: model.feature_importances_
```

Out[102]: array([0.17932645, 0.03142367, 0.20624864, 0.02317516, 0.02318449,
0.01660314, 0.05738771, 0.03673533, 0.19037076, 0.15401823,
0.08152642])

```
In [103]: ranked_features = pd.Series(model.feature_importances_, index = A.columns)
```

```
In [104]: ranked_features.nlargest(12).plot(kind='bar') # rank and plot the scores
```

Out[104]: <AxesSubplot:>



```
In [105]: mu_ifo = mutual_info_classif(A,b)
mu_data = pd.Series(mu_ifo, index = A.columns)
mu_data.sort_values(ascending=False)
```

```
Out[105]: age                0.034894
work_type            0.011995
heart_disease        0.010218
bmi                  0.008821
smoking_status       0.007195
id                   0.005207
avg_glucose_level    0.005163
hypertension         0.004164
ever_married         0.001049
Residence_type       0.001045
gender               0.000000
dtype: float64
```

```
In [106]: a_train,a_test,b_train,b_test = train_test_split(A,b, test_size = 0.2, random_state=42)
```

```
In [107]: def classify(model, a, b):
            a_train,x_test,b_train,b_test = train_test_split(A,b, test_size = 0.2, random
            model.fit(a_train,b_train)
            print('Accuracy is: ', model.score(a_test,b_test)*100)
            score = cross_val_score(model,a,b,cv=5) # we set the number
            print('Cross validation Accuaracy: ', np.mean(score)*100)
```

```
In [108]: LG_model = LogisticRegression()
            classify(LG_model, A,b)


Accuracy is: 93.9334637964775
Cross validation Accuaracy: 95.06849315068493
```

```
In [109]: b_pred = LG_model.predict(a_test)
```

```
In [110]: com = confusion_matrix(b_test,b_pred)
```

```
In [111]: sns.heatmap(com,annot=True) # plots the confusion matrix
```

Out[111]: <AxesSubplot:>



```
In [112]: knn_model = KNeighborsClassifier(n_neighbors=5)
            classify(knn_model,A,b)
```


```
Accuracy is: 93.9334637964775
Cross validation Accuaracy: 95.12720156555773
```

```
In [113]: b_pred = knn_model.predict(a_test)
```

```
In [114]: com = confusion_matrix(b_test,b_pred)
```

```
In [115]: sns.heatmap(com,annot=True) # plots the confusion matrix
```

Out[115]: <AxesSubplot:>



```
In [116]: SVC_model = SVC(kernel='linear', C = 1)
          : classify(SVC_model, A,b)
```

Accuracy is: 93.34637964774952

Cross validation Accuracy: 94.28571428571428

```
In [117]: SVC_model = SVC(kernel='linear', C = 1)
```

```
In [118]: SVC_model.fit(a_train,b_train)
```


Out[118]: SVC(C=1, kernel='linear')

```
In [119]: b_pred = SVC_model.predict(a_test) # save predictions in y_pred
```

```
In [120]: com = confusion_matrix(b_test,b_pred)
```

```
In [121]: sns.heatmap(com,annot=True) # plots the confusion matrix
```

```
Out[121]: <AxesSubplot:>
```



THANK YOU

```
In [ ]:
```